

MercadoLibre Inc. (MELI.US)

Dominant LatAm marketplace + payments + credit + logistics platform, trading near intrinsic value, with a structural overlay on AI-disclosure quality.

Sector: Consumer Cyclical · Internet Retail · Broadline Retail
Listing: NASDAQ · IPO 2007-08-10 · HQ Montevideo, Uruguay
Reference price: ~\$1,648 · **Market cap:** ~\$83.55bn · **EV:** ~\$90.44bn
As of: 2026-05-28

STARTER POSITION · CONDITIONAL

§0 — INVESTMENT GRADE VERDICT

STARTER POSITION (CONDITIONAL)

Position size: 0.5 – 1.0% of NAV (~\$5m–\$10m per \$1bn NAV)
Entry approach: small initial bite at spot; scale on retracement to the \$1,270 – \$1,430 band (10–20% margin of safety to base case)
Conviction: Thesis HIGH · Price FAIR-LEANING-EXPENSIVE · Information sufficiency MEDIUM

"Dominant LatAm marketplace + payments + credit + logistics + advertising platform compounding at industry-leading rates, trading near base-case intrinsic value, with a structural bear-case overlay on the quality of AI-disclosure underlying the margin-expansion path."

STRATEGY VERDICT — LOCKED

RECOMMENDATION	INITIAL \$	MARGIN OF SAFETY	FULL AT
----------------	------------	------------------	---------

STARTER**\$50,000****-1.0%****≤ \$1,270**

MercadoLibre is the integrated marketplace + payments + credit + logistics + advertising platform of Latin America, compounding revenue at 42.2% over the four-year window 2021-2025 and re-accelerating to +49% YoY in Q1 2026 [workstream_d Section 2]. The wide moat is structural (Helmer 7 Powers weighted 115.7; scale economies, network economies, switching costs all load) rather than AI-exceptionalist [workstream_b moat-assessment]. At spot \$1,648 the probability-weighted intrinsic is \$1,632, implying -1.0% margin of safety [workstream_d Section 7]. The verdict is STARTER, not full, for two reasons. First, the bull case (which re-rates intrinsic to ~\$1,925) requires AI-cohort disclosure with attribution that does not yet anchor to primary management filings [cross_red_flags RF-DISC-006]. Second, the price embeds consensus 17-19% revenue CAGR at a developed-market discount rate — the moat is real but the moat is not paying for itself in margin-of-safety terms at spot. Take a starter seat sized at \$50k pending Szarfsztein's first AI-cohort disclosure; full position only at \$1,270 or below (20% MoS). The asymmetric upside is genuine; the patience cost is small.

+49.0%

Q1'26 YOY REVENUE

42.2%

4-YR REVENUE CAGR

43.5x

TRAILING P/E

\$1,632INTRINSIC (PROB-
WTD)**-1.0%**

MOS AT SPOT

\$12.1bn

NET CASH

131

MOAT (VS 100 MED)

9/86

RED FLAGS FIRED

FOUR CONDITIONAL TRIGGERS TO UPSIZE FROM STARTER TO FULL

1. **AI cohort metrics anchored to primary disclosure.** Szarfsztein's Q2 or Q3 2026 earnings call — or the FY2026 10-K MD&A — names at least one specific AI-driven cohort metric with attribution. Discharges RF-DISC-006 and re-rates the bull-case probability from 25% to ~40%, moving probability-weighted intrinsic from \$1,632 to approximately \$1,925. [T1 — cross_red_flags \$RF-DISC-006]
2. **Brazil regulatory clarity.** BCB publishes a rule that draws a stable perimeter around platform-originated consumer credit, or formally indicates a 12+ month moratorium on further tightening. [T1 — workstream_b \$regulatory]

3. **Q1 2026 +49% YoY confirmed as the new run-rate, not an anomaly.** Q2 and Q3 2026 print at 35%+ YoY revenue growth on clean comps. The implied five-year revenue compound moves from the consensus 17–19% band toward the 25–30% band. [T1 — *workstream_d* §scenario-probability-weights]

4. **Argentina FX stability.** Sustained peso stabilisation through a full calendar quarter, or MELI moving to constant-currency primary reporting comparable across periods. Removes IAS 29 hyperinflationary noise from headline GAAP.

[T1 — *workstream_e/primary-filings*]

§1 — EXECUTIVE SUMMARY

What MercadoLibre is. The integrated commerce-payments-credit-logistics-advertising platform of Latin America. Founded in Buenos Aires in 1999 by Marcos Galperín and listed on NASDAQ in 2007. The company has compounded revenue at a 42% four-year CAGR through FY2025 [T1 — *workstream_d* §2], with a Q1 2026 print re-accelerating to +49% YoY — the strongest quarter since Q2 2022 [T1 — *workstream_d* §2; 10-Q 2026-05-08]. The platform operates across 18 LatAm countries and owns first-party logistics through Mercado Envíos, licensed fintech operations through Mercado Pago and Mercado Crédito, and a fast-growing advertising arm in Mercado Ads [T1 — *workstream_b/competitor-analysis* §moat].

The thesis is structural, not AI-exceptionalist. MercadoLibre clears Mauboussin's "wide moat" bar on the visible structural ecosystem — scale economies, network economies, and switching costs across the five-asset stack — not on AI exceptionalism [T1 — *workstream_b* §moat]. The moat-durability index runs 131 versus a sector median of 100 [T1 — *workstream_d* §5]. Twenty-six years of accumulated thin-file LatAm behavioural data and a multi-billion-dollar logistics build-out across BR/MX/CL/CO are the durable physical and informational moats; the Helmer 7 Powers weighted score reads 115.7 with the moat index at 125 [T1 — *workstream_b* §moat].

The verdict is starter, not full, because the price is fair-leaning-expensive. At a reference price of \$1,648 (market cap \$83.55bn divided by 50.697m diluted shares; price identity verified through \$37.90 TTM EPS × 43.48 P/E) [T1 — *workstream_d* §1], MercadoLibre trades at a probability-weighted intrinsic of \$1,632 — a margin of safety of –1.0% [T1 — *workstream_d* §7]. The 4-year revenue compound is exceptional and the Q1'26 reacceleration is real. But the price embeds a consensus 17–18% revenue CAGR over

five years at a developed-market discount rate, and the AI narrative the company invokes to defend the margin-expansion path does not yet anchor to audited cohort disclosure.

BEAR CASE — THE OPERATIVE RISK, NOT A TAIL OUTLIER

Workstream A grades zero of MercadoLibre's six AI deployments at evidence grade B or better. The three load-bearing AI claims the bull thesis relies on — Mercado Pago's 87% AI customer-service resolution rate, Mercado Crédito's 23.3% NIMAL, and the 40% YoY fraud-cost decline — are sourced from secondary synthesis (klover.ai, investing.com, Gate 6 reference), not primary management disclosure [T1 — cross_red_flags §RF-DISC-006]. This fires the one high-severity flag in a 9-flag sweep across an 86-flag canonical catalogue. It is the canonical Schilit "reliance on adjusted financial metrics" pattern, inverted into modern AI-narrative form.

The position-sizing call is explicitly conditional. Take a starter position today at 0.5–1.0% of NAV. Use the next two earnings prints — Q2 and Q3 2026 — as the deciding evidence. If Szarfsztejn (five months into his tenure as the first non-founder CEO) names a specific AI cohort metric with attribution, the bull-case probability re-rates from 25% to ~40%, the probability-weighted intrinsic moves from \$1,632 to roughly \$1,925, and the starter scales to a full position. If he defaults to a scale-and-reinvestment narrative for a fourth consecutive reporting year, the base-case probability weight should compress toward the bear scenario (\$623/share) and the position is held flat or trimmed.

What would change our mind. Three falsifiability triggers operate concurrently. First, a single quarter in which Mercado Crédito 90+ day non-performing loans expand by more than 200 basis points without a quantitative explanation. Second, a Shopee LatAm seller-subsidy programme that compresses take-rate from 20–22% to 17–18%. Third, a BCB rule that caps platform-originated consumer credit at a sub-Mercado-Crédito-current-rate growth level. Any one materially shifts the bear-case probability mass.

Indexes: Moat durability 131 vs 100 median. Helmer 7 Powers weighted 115.7. AI disclosure quality ~70 vs 100 median.

§2 — COMPANY OVERVIEW

From a Buenos Aires marketplace to an 18-country integrated platform

DATE	EVENT	SOURCE
1999	Marcos Galperín co-founds MercadoLibre in Buenos Aires as a regional online marketplace	[T1 — <i>workstream_a §events</i> ; 10-K FY2025]
2004	Mercado Pago payments launched under Osvaldo Giménez	[T1 — <i>workstream_a §events</i>]
2007-08-10	NASDAQ IPO — first major LatAm tech listing on a US exchange	[T1 — <i>workstream_a §events</i> ; SEC 424B1]
2017–2018	Mercado Crédito scaled under de los Santos; Mercado Envíos in-house logistics build-out begins under Szarfsztein	[T1 — <i>workstream_a §events</i>]
2026-01-01	Galperín transitions to Executive Chairman; Szarfsztein becomes CEO — first non-founder CEO in company history	[T1 — <i>workstream_a §events</i> ; DEF 14A]

Five complementary businesses on one ecosystem

MercadoLibre is structurally five businesses on one ecosystem. **Mercado Libre Marketplace** is the two-sided commerce platform across 18 LatAm countries, with a take rate of 20–22% [T1 — *workstream_b §MC-01*; 10-K]. **Mercado Pago** is the digital payments and fintech-services platform. **Mercado Crédito** is the consumer and seller credit book, built on platform behavioural data. **Mercado Envíos** is the first-party logistics network across the five priority markets; unit shipping costs fell YoY in Brazil, Mexico, Chile, and Colombia in 2025 per the Q4 release, and management explicitly attributes the decline to scale, not AI [T1 — *workstream_a §dep-004*; BusinessWire 2026-02-24]. **Mercado Ads** is the advertising platform; ad revenue runs at ~2% of GMV and is growing 50–67% YoY

[T3 — *workstream_a §dep-003*].

Headline corporate facts

- **Incorporation:** Delaware. **HQ:** WTC Free Zone, Montevideo, Uruguay

[T1 — EODHD General]

- **CIK:** 0001099590 · **CUSIP:** 58733R102 · **ISIN:** US58733R1023 [T1 — EODHD]

- **Auditor:** not identified in current corpus — Gate 8 confirmation item

[T1 — cross_red_flags \$RF-AUD-001]

- **Diluted shares outstanding:** 50,697,182 [T1 — ipo_basics.json; EODHD SharesStats]

- **Full-time headcount (FY2025):** 123,670 [T1 — workstream_a \$events; 10-K]

Indexes: Operating perimeter 18 countries. Headcount 123,670. Diluted shares 50.697m.

§3 — MARKET OPPORTUNITY

The runway is the penetration gap, not the TAM number. LatAm e-commerce penetration sits below 15% across the largest markets, against a developed-market range above 20% — the structural runway is the gap between today and the developed-market range [T2 — workstream_b \$substitutes]. The penetration gap is the load-bearing argument. Closing half of it over a decade is consistent with the consensus 17–19% revenue CAGR the current price embeds in the reverse-DCF [T1 — workstream_d \$reverse-DCF].

Market structure across the three legs

- **Commerce GMV.** The combined Brazil-Mexico-Argentina commerce TAM is the headline, with secondary markets (Chile, Colombia, Peru, Uruguay) layering on. Brazil is MercadoLibre's largest single revenue contributor at an estimated ~55% of total [inference — workstream_d \$WACC build].
- **Payments and fintech.** Mercado Pago competes with bank cards, Nubank's neobank stack, and traditional payment processors. LatAm consumer credit cards have low thin-file penetration, which is the structural opening for Mercado Crédito [T2 — workstream_b \$direct-competitor].
- **Logistics.** Mercado Envíos competes with regional carriers (Correios in Brazil, MexPost) and global parcel networks (DHL, FedEx) at the long-haul end. First-party logistics across 18 countries is a multi-billion-dollar capex moat that no LatAm peer has matched [T1 — workstream_b \$EM-native-moat].

The reference class — Amazon, Sea, MELI Q1'26

The market opportunity is not a TAM-multiple exercise. It is the thesis that the integrated five-asset stack monetises a single customer relationship more deeply than any single-vertical competitor can. The four-year revenue compound of 42.2%

[T1 — *workstream_d §2*] argues the cross-asset compounding is real; the Q1 2026 +49% YoY print [T1 — *workstream_d §2*] argues it has not yet peaked.

The reference class for forward growth is LatAm-listed integrated platforms over rolling five-year windows. Sea Limited (Shopee), the structural Asia analogue, posted 38.4% TTM revenue growth [T2 — *workstream_d §3*]. Amazon, the global integrated-commerce reference, grew 13.6% TTM. MELI's Q1 2026 +49% is above both. The reference-class base rate is therefore a wide band: 14% (Amazon) at the mature end, 38% (Sea) at the comparable end, 49% (MELI Q1) at the recent print. Consensus models 17–19% revenue CAGR over five years [T2 — *workstream_g §bull-framing*]; our base case adopts a similar trajectory.

Indexes: LatAm e-commerce penetration ~75 vs ~100 developed-market median. MELI Q1'26 growth vs Sea reference class ~128.

§4 — COMPETITIVE LANDSCAPE

Three direct threats — Shopee, Amazon LatAm, Nubank

Shopee (Sea Limited) is the most-watched threat. A marketplace + payments + credit stack expanding into LatAm with SE-Asia cashflows behind it

[T1 — *workstream_b §direct-competitor*]. The war-game scenario — Shopee accelerates LatAm seller subsidies funded by SE-Asia cashflows while Nubank expands Brazil/Mexico consumer credit at lower rates than Mercado Crédito — is the central bear-case mechanism in the valuation [T1 — *workstream_b §war-game*]. It is the most concrete take-rate-compression path from 20–22% to 17–18%, the band that drives MELI's bear-case intrinsic of \$623/share [T1 — *workstream_d §6*].

Amazon LatAm competes head-on in commerce, logistics, and advertising but has historically been sub-scale in LatAm payments and credit [T2 — *workstream_b §direct-competitor*].

Amazon's mature US ad-revenue-to-GMV ratio runs 7–8% [T2 — *workstream_a \$dep-003*], versus MELI's ~2% — meaningful headroom, but Amazon must execute LatAm-specific logistics density to capture it.

Nubank attacks the fintech arm directly. Nubank's AI-native credit underwriting across Brazil, Mexico, and Colombia is the structural competitor to Mercado Crédito on cost-of-funds and credit-decision velocity [T2 — *workstream_b \$direct-competitor*]. Where Mercado Crédito earns its margin on platform-data-enabled thin-file underwriting, Nubank earns its margin on regulated-bank cost-of-funds advantage. The two coexist today; the question is whether Nubank's bank-rate funding advantage erodes Mercado Crédito's NIMAL over time.

MercadoLibre's moat — three powers load, two do not

The Helmer 7 Powers read scores MELI as a wide-moat candidate, weighted 115.7, with a moat index of 125 [T1 — *workstream_b \$moat*]. Three powers load:

- **Scale economies (135).** Mercado Envíos unit shipping costs fell YoY in Brazil, Mexico, Chile, and Colombia per the Q4 2025 release [T1 — *workstream_b \$moat*; *BusinessWire 2026-02-24*]. Management explicitly attributes the decline to scale.
- **Network economies (135).** Eighteen-country two-sided marketplace liquidity is primary-sourced to 10-K Item 1 [T1 — *workstream_b \$moat*].
- **Switching costs (130).** Sellers depend on Mercado Pago for settlement, Mercado Envíos for fulfilment, and Mercado Crédito for working capital simultaneously [T1 — *workstream_b \$moat*; 10-K].

Process power from the GenAI Gateway on the "Fury" platform scores lower at 105 because the model layer is increasingly commodity [T2 — *workstream_b \$moat*]. Cornered resource scores 95 — 26 years of accumulated thin-file LatAm behavioural data is replicable with time, not legally protected.

Voice-of-customer overlay

STATUS: FLAG-ONLY, NOT LOAD-BEARING

The Voice-of-Customer agent did not retrieve a primary post-level corpus in this run, and the doctrine §4 gate-1 floor (N=8 floor with citations) is not cleared [T4 — *workstream_b/voice-of-customer §run-level-finding*]. Four candidate themes appear as colour signal only: (i) take rate at the pain threshold (seller discourse frames combined marketplace + Envíos + Pago fees as already margin-crushing); (ii) trust-and-safety dispersion in electronics/apparel/cosmetics (the categories MELI most needs to defend); (iii) Mercado Pago friction (account locks, KYC freezes, dispute delays); (iv) logistics expectations as a liability (FULL branding raises user expectations to Amazon Prime levels). None promoted to load-bearing finding.

Regulatory overlay by jurisdiction

JURISDICTION	REGULATOR STACK	TIGHTENING HALF-LIFE
Brazil	BCB + LGPD	12–24 months on fintech-credit origination [T1 — <i>workstream_b §regulatory</i>]
Argentina	BCRA + CNV	6–12 months on FX controls and capital repatriation [T1 — <i>workstream_b §regulatory</i>]
Mexico	CNBV + Banxico	12–24 months on Ley Fintech compliance [T1 — <i>workstream_b §regulatory</i>]
USA	SEC	Stable; listed-issuer disclosure standard [T1 — <i>workstream_b §regulatory</i>]

Indexes: Moat structure 125. Scale / network / switching all 130–135. Process power 105. Cornered resource 95.

§5 — MANAGEMENT TEAM

A Wasserman-textbook founder transition, executed before pressure required it

Marcos Eduardo Galperín co-founded MercadoLibre in 1999 in Buenos Aires and ran it as CEO for more than 26 years, through Stanford MBA, the 2007 IPO, scale-up across

18 LatAm markets, and the build-out of the Mercado Pago, Mercado Crédito, Mercado Envíos, and Mercado Ads adjacencies [T2 — *workstream_c §founder-transition*]. He announced on 21 May 2025 that he would step down as CEO effective 31 December 2025 and assume the Executive Chairman role from 1 January 2026, with Ariel Szarfsztejn promoted from Commerce President to CEO and President on the same date

[T2 — *workstream_c*; corroborated DEF 14A 2026-04-23].

Galperín retains chairmanship of the board, public-facing strategy custody, capital-allocation authority, and AI-innovation sponsorship per the transition announcement [T3 — *workstream_c §founder-transition*]. This is a Wasserman-textbook founder-role-sacrifice transition (Marker D), executed before pressure required it: governance, succession, and equity-control questions handed off voluntarily by a founder still well within operating prime. The founder-courage marker scores ~110 on the 100-equals-peer-median index

[inference — *workstream_c*].

Szarfsztejn — what we know after five months in the seat

Ariel Szarfsztejn joined MercadoLibre in 2017 as VP Strategy and Corporate Development, ran Mercado Envíos logistics for approximately four years (2018–2021), then served as EVP Commerce from January 2022 and Commerce President from January 2024 through December 2025 [T2 — *workstream_c*]. Pre-MELI he was at Despegar (LatAm online travel) and Boston Consulting Group [T2]. Education: economics degree (cum laude) from Universidad de Buenos Aires; MBA from Stanford GSB [T2].

Total annual compensation aggregates around US\$4.3m, of which approximately 15% is salary and 85% is variable or stock-based [T3 — *workstream_c §compensation*]. Direct equity ownership is very small at approximately 0.0001% of shares outstanding [T3 — *workstream_c*]. The compensation structure is conventional listed-equity professional-CEO design, heavily skewed to stock-linked variable pay — alignment is via grant value, not founder-level equity.

The track-record peer index reads ~120 absolute / ~112 shrunk-peer at medium confidence [inference — *workstream_c*]. The logistics build-out is independently corroborated as a moat driver [T3 — *Quart*], but the CEO seat itself has zero quarterly delivery on the record. He is five months in seat at the date of this paper.

Bench depth — Rabinovich, de los Santos, Giménez carry the operating stack

Daniel Rabinovich — COO (Product and Technology) — is the single most load-bearing operational figure after the founder. He joined MELI as an application architect in March 2000, became VP Product Development in 2009, CTO in January 2011, COO Product and Technology in 2019, and full COO in August 2020 [T2 — *workstream_c* §bench]. Twenty-six continuous years at the platform, owning the engineering and product stack across the AI deployment effort. No prior AI-lab affiliation [T2].

Martín de los Santos — EVP and CFO since January 2024 — joined MELI in 2013 as VP Strategy, then led Mercado Crédito from 2017 to 2023. Pre-MELI: Vostu, IMPSA, Merrill Lynch, McKinsey, Goldman Sachs. Stanford MBA, UNC Chapel Hill BS [T2 — *workstream_c*]. Running Crédito put him at the centre of the data-and-risk-modelling stack that underwrites the fintech monetisation thesis.

Oswaldo Giménez — Fintech President (Mercado Pago) — joined MELI in 2000 and has run Mercado Pago since 2004. Pre-MELI: Booz Allen Hamilton and Santander Investments. Industrial engineering degree, Stanford MBA [T2 — *workstream_c*].

Other named EVPs include Agustín Costa (Shipping), Fernando Yunes (EVP and Head of Commerce LatAm / Brazil lead), Sean Summers (CMO and EVP Advertising), and Juan Martín de la Serna (Corporate Affairs and President Argentina) [T2 — *workstream_c*]. Median shared tenure across the named team is unusually high — at least four of the seven named EVPs have been at MELI for more than 10 years.

Key-person risk concentrates at Galperín and Rabinovich

Two simultaneous concentration points. First, Galperín as Executive Chair plus founder gravity: if he withdraws fully in the next 24 months, the cultural and strategic-continuity bridge to Szarfsztejn shortens fast. Second, Rabinovich as the 26-year owner of product and technology: single-point-of-failure across the engineering stack supporting the AI deployment story. Bench is deep but no publicly-named successor to either is visible. Key-person risk: medium [T2 — *workstream_c* §key-person-risk].

Culture and AI readiness

Glassdoor — reported as substantial Spanish-language review volume — shows 84% recommend-to-friend and 3.6/5 work-life-balance rating, with recurring praise for collaborative environment and dissenting reviews citing harsh internal communication and a challenging pace [T3 — *workstream_c §culture*]. Trailing-12-month delta not extractable from current sources [unverified].

Szarfsztein has publicly described running centralised AI-training programs for senior leaders including himself and Galperín, and is deploying agentic AI features in Mercado Pago and the marketplace app [T3 — *workstream_c; AS/COA 2024 talk*]. No headline hires from OpenAI, Anthropic, DeepMind, or peer AI labs into C-suite or named EVP roles are evident in the sources [T2 — *workstream_c*]. No board director carries a primary AI-research or AI-engineering identity — Dubugras (Brex) is the closest fintech-tech founder profile, and Tolda brings deep e-commerce operating depth, but AI sits entirely inside management [T2 — *workstream_c §board*].

PEOPLE & CULTURE — IC DISCUSSION FLAGS

(1) The CEO has zero full quarters of public results in seat. Underwrite the next two earnings cycles as the first real evidence on Szarfsztein versus the founder-era operating trajectory.

(2) The board carries no director whose primary identity is AI research or AI engineering. If the AI-flywheel thesis is the entry justification, consider whether a director-level AI seat is a precondition to upsizing from starter to full.

Indexes: Founder-courage marker ~110. Szarfsztein track-record ~120 absolute / ~112 shrunk-peer. Bench depth and median EVP tenure unusually high.

§6 — FINANCIAL SUMMARY

Five-year revenue, margin, and earnings trajectory

YEAR	REVENUE (\$M)	YOY	OP. INCOME (\$M)	EBITDA (\$M)	NET INCOME (\$M)	OP. MARGIN
FY2021	7,069	—	441	674	83	6.24%
FY2022	10,780	+52.5%	1,069	1,318	482	9.92%
FY2023	15,107	+40.1%	2,207	2,251	987	14.61%
FY2024	20,777	+37.5%	2,631	3,214	1,911	12.66%
FY2025	28,893	+39.1%	3,201	3,561	1,997	11.08%

Source: [T1 — workstream_d §2; EODHD Income Statement; MELI 10-K filings 2026-02-25 for FY25]

The compound numbers. Four-year revenue CAGR 2021–2025 = 42.2%; three-year CAGR 2022–2025 = 38.9%. The Q1 2026 quarterly print was \$8,845m versus \$5,935m in Q1 2025 — a +49.0% YoY print and the strongest YoY revenue print since Q2 2022

[T1 — workstream_d §2; 10-Q 2026-05-08].

Margin compression is the central debate. Operating margin compressed in FY25 to 11.08% (down from 12.66% in FY24) as the company reinvested fintech credit-book growth into provisions and shipping subsidies [T1 — workstream_d §2]. TTM operating margin through Q1 2026 reads 6.91% [T1 — workstream_d §1].

Headline TTM ratios

EBITDA margin **12.39%**; gross margin **49.47%**; profit margin **6.04%**; return on equity **31.26%**; return on assets **5.11%**. Diluted EPS TTM **\$37.90**; FY26E EPS estimate **\$41.08** (forward P/E 31.95x); FY27E EPS estimate **\$58.63** [T1 — EODHD Highlights / Valuation].

Balance sheet — \$12.1bn net cash carries the position

Q1 2026 balance sheet shows total interest-bearing debt of \$5.0bn against \$17.1bn in cash and short-term investments — **net cash of \$12.1bn at 31-Mar-2026**

[T1 — workstream_d §1; 10-Q 2026-03-31]. Reported enterprise value of \$90.4bn carries a \$6.9bn premium to market cap that reflects the Mercado Pago funding and deposits liability footprint, not net financial debt. This wedge is the most important reconciliation item in the financial model: EV-based multiples understate "cleanness" relative to the equity-bridge view.

Cash flow — the float caveat

TTM free cash flow of \$10.71bn (33.7% FCF margin) is materially inflated by Mercado Pago float dynamics — customer deposit accumulation runs through operating cash flow under GAAP *[T1 — workstream_d §1]*. A normalised operating-FCF view stripping float would land closer to **\$3–4bn** — meaningfully thinner than the headline 12.8% FCF yield suggests *[inference — workstream_d §1]*.

Earnings cadence and management posture

Management language across the last four reportable quarters projects disciplined bullishness: 30%+ revenue growth, record operating income in Q2 2025, and an improving credit-quality narrative *[T1 — workstream_e/earnings-transcripts]*. CFO Martín de los Santos has been the consistent voice on the calls. Incoming-CEO Szarfsztein does not appear in available Q2/Q3 2025 excerpts; his Q1 2026 earnings-call commentary is the first opportunity to read his quarterly posture.

The clearest, most quantitatively anchored language came in Q2 2025: **ads revenue +38% YoY, 50–90 day NPLs below 7% for the first time** since reporting began, **NIMAL broadly stable QoQ**, and more than half of the Brazil credit book NIMAL-positive *[partial-T1 — workstream_e §verbatim-signal; Q2-25 10-Q anchors pending]*. Mexico GMV "accelerated sharply" with items sold at the "fastest pace in almost 2 years." Unambiguous, superlative-laden statements — not hedged.

Tone read: BULLISH. No defensive posture, no walk-back, no apology for investment levels. MELI does not issue numerical guidance *[T1 — workstream_e §tone]*; the cultural pattern is qualitative-only forward commentary. Adversarial flags from the transcript: (i) AI attribution gap — management invokes ML in credit, ads, search, and logistics but no verbatim quote isolates AI's contribution to margin or take-rate; (ii) peak-cycle superlatives stacked in one quarter warrants watching whether Q4 2025 sustains (Q4 +45% YoY suggests yes); (iii) CEO-transition silence — Szarfsztein does not appear in surfaced excerpts; (iv) no downside-scenario credit language — only positive asset-quality framing, itself a tell in a high-growth LatAm credit book.

Investor sentiment overlay

Investor sentiment is **mixed and notably more sober than sell-side consensus**

[T2 — workstream_h §bottom-line] — a cleaner setup than the typical retail-euphoria pattern. There is no named short-seller report (Hindenburg, Muddy Waters, Spruce Point, Citron — all silent on MELI) *[T2 — workstream_h]*. Short interest is reportedly low and declining *[T2; unverified]*. This is not a crowded short, not a meme name, not a gamma-squeeze candidate.

The dominant retail framing is "*great business, possibly fully valued, credit risk underappreciated*" — not the rocket-emoji pattern that typically functions as a contrarian sell signal. Retail being less bullish than analysts on a quality compounder is mildly constructive. Counter-evidence on the AI narrative: implicit skepticism manifesting as "if AI is so transformative, where is the margin payoff?" — the same disclosure-quality concern that fires the high-severity red flag at the workstream level.

Indexes: TTM ROE 31.26%. Gross margin 49.47%. Net cash position \$12.1bn. Headline FCF flatters by roughly three to one.

§7 — VALUATION & ENTRY PRICE

Current market position — 38% below 12-month high

MELI trades at an implied **\$1,648/share** (market cap \$83.55bn divided by 50.697m diluted shares; identity verified: \$37.90 TTM EPS × 43.48 trailing P/E = \$1,648 to two decimals) *[T1 — workstream_d §1; EODHD highlights 2026-05-28]*. The 52-week range is **\$1,495–\$2,645**, with the 50-DMA at \$1,726 and the 200-DMA at \$2,036; the stock is presently trading **19% below its 200-DMA and 38% below its 12-month high** — a meaningful correction from a peak Mag-7-style multiple *[T2 — workstream_d §1; EODHD Technicals]*.

MULTIPLE (TTM)	MELI	SOURCE
Trailing P/E	43.48x	<i>[T1 — EODHD]</i>
Forward P/E (FY26E EPS \$41.08)	31.95x	<i>[T1 — EODHD]</i>
EV / Revenue (TTM \$31.80bn)	2.84x	<i>[T1 — EODHD]</i>
EV / EBITDA (TTM \$3.94bn)	23.84x	<i>[T1 — EODHD]</i>

MULTIPLE (TTM)	MELI	SOURCE
Price / Sales (TTM)	2.63x	[T1 — EODHD]
Price / Book (MRQ)	11.48x	[T1 — EODHD]
TTM FCF yield (float-inflated)	12.82%	[T1 — EODHD; float caveat]
PEG	1.02x	[T1 — EODHD]

Peer comparables — between Sea and Amazon on EV/Revenue

TICKER	NAME	EV/REV	EV/EBITDA	P/E	REV GTH	WHY-COMPARABLE
MELI.US	MercadoLibre	2.84x	23.84x	43.48x	+49.0%	Anchor
AMZN.US	Amazon	3.8x	16.5x	34.9x	+13.6%	Integrated commerce + logistics + payments + ads
SE	Sea Limited (Shopee)	2.0x	20.2x	36.2x	+38.4%	Closest structural analogue
PYPL.US	PayPal	1.4x	5.9x	9.4x	+3.7%	Payments-only; low-growth incumbent
—	Mendoza peer median	3.9x	14.6x	34.9x	n/a	Mendoza Feb 2026 equity report

Sources: [T1 — EODHD for MELI; T2 — Gate 6 baseline EODHD-API for peers; T3 — Mendoza report]

Peer-relative read. On EV/Revenue, MELI's 2.84x sits between Sea (2.0x) and Amazon (3.8x). On EV/EBITDA, MELI's 23.84x is a 17% premium to Sea and a 44% premium to Amazon — premium-priced but not extreme. On trailing P/E, the 43.48x is 25% above Amazon and 4.6x PayPal. **Verdict: fair, leaning stretched.** The premium is defensible only if the 49% Q1 print is sustained beyond a single quarter [T1 — *workstream_d \$3*].

DCF — Chronos DCF Standard v1 build

WACC

COMPONENT	VALUE	SOURCE / RATIONALE
Risk-free rate (10y UST)	4.60%	DCF Standard \$4 default <i>[T0]</i>
Mature-market ERP	4.21%	Damodaran July 2025 <i>[T0]</i>
Revenue-weighted CRP	5.15%	Brazil 55% × 3.0% + Argentina 20% × 14.0% + Mexico 18% × 3.1% + Other 7% × 2.0% <i>[T0; inference]</i>
Beta (vs S&P, 2y weekly)	1.41	<i>[T1 — EODHD Technicals]</i>
Cost of equity (Damodaran melded)	17.76%	RFR + $\beta \times$ (ERP + CRP)
Cost of equity (McKinsey-lite)	~10.5%	McKinsey 2021 critique of melded CRP <i>[T0]</i>
House blended cost of equity	13.5%	Bracketed average — defensible for US-listed USD-reporting LatAm operator
Pre-tax cost of debt	7.0%	LatAm fintech blended USD facility <i>[inference]</i>
After-tax cost of debt	5.25%	25% effective tax
Capital weights (E / D)	94.4% / 5.6%	Market values
WACC (base)	13.5%	House blended

FIVE-YEAR EXPLICIT FORECAST — BASE CASE

YEAR	REV (\$M)	YOY	OP MARGIN	EBIT	NOPAT	D&A	CAPEX	ΔNWC	UFCF
FY2026	36,116	+25%	10.9%	3,929	2,947	903	1,625	289	1,936
FY2027	44,062	+22%	12.2%	5,358	4,018	1,102	1,983	318	2,819
FY2028	52,434	+19%	13.4%	7,047	5,285	1,311	2,360	335	3,902
FY2029	61,347	+17%	14.7%	9,030	6,773	1,534	2,761	357	5,189

YEAR	REV (\$M)	YOY	OP MARGIN	EBIT	NOPAT	D&A	CAPEX	ΔNWC	UFCF
FY2030	71,163	+16%	16.0%	11,386	8,540	1,779	3,202	393	6,724

Followed by a 5-year fade FY2031–2035 with growth linearly converging from 16% → 3% and operating margin expanding to 18% (Amazon's run-rate), then Gordon perpetuity at 3.0% terminal growth. Explicit + fade + Gordon, all three TV methods reconciled within ±25% per DCF Standard §5.

DCF RESULTS — THREE SCENARIOS

SCENARIO	WACC	5Y GROWTH	TERM OP MGN	TERM GTH	EV (\$M)	TV % EV	EQUITY (\$M)	PRICE/SH
Bull	13.0%	40/30/25/20/18%	22%	3.5%	126,079	57%	138,182	\$2,726
Base	13.5%	25/22/19/17/16%	18%	3.0%	68,456	53%	80,559	\$1,589
Bear	15.0%	15/12/10/8/7%	10%	2.5%	19,504	39%	31,607	\$623

All three TV-as-%-of-EV figures sit comfortably below the 75% amber threshold and the 85% red threshold per DCF Standard §5 — the model is structurally healthy.

SCENARIO PROBABILITY WEIGHTS AND 5-YEAR IRR

SCENARIO	PROBABILITY	PRICE/SH	5Y IRR FROM \$1,648	ANCHOR FOR WEIGHT
Bull	25%	\$2,726	+10.6%	Q1'26 +49% sustains; AI-credit flywheel compounds; Shopee retreats
Base	50%	\$1,589	−0.7%	Sell-side convergence at ~17% revenue CAGR + margin expansion to 18%
Bear	25%	\$623	−17.7%	Shopee LatAm subsidy → take-rate compression; BCB tightens; execution gap

PROBABILITY-WEIGHTED INTRINSIC AND MARGIN OF SAFETY

Probability-weighted intrinsic value: \$1,632/share. Versus current \$1,648 → **margin of safety = −1.0% (no margin of safety at spot).** Bull-bear IRR spread = 28.3 pp — narrowly below the DCF Standard \$6 30 pp floor. Justification: LTM revenue base \$31.8bn; outcome tails are narrower than the Standard's early-stage default.

REVERSE-DCF

At WACC = 13.5% the current \$1,648 price embeds a **5-year revenue CAGR of approximately 18%** with terminal operating margin reaching 17% — broadly consistent with sell-side consensus (17% revenue CAGR, 24% EPS CAGR over next 3 years per Simply Wall St aggregation) *[T3 — workstream_d \$reverse-DCF]*. At WACC = 12% the same price embeds 15% revenue CAGR — undemanding. At WACC = 15% the price embeds >25% revenue CAGR — demanding. **The market is pricing consensus growth at a developed-market discount rate.** That is the central debate.

MULTIPLES CROSS-CHECK (FY26E BASE-CASE ANCHORS)

- **EV / Sales:** FY26E revenue \$36.1bn × 2.0x (Sea multiple) = \$1,664/sh; × 2.5x = \$2,020/sh; × 3.0x = \$2,376/sh
- **EV / EBITDA:** FY26E EBITDA ≈ \$5.96bn × 18x = \$2,355/sh; × 20x = \$2,590/sh
- **P/E (FY26E):** EPS \$41.08 × 40x = \$1,643/sh (≈ current); × 30x = \$1,232/sh; × 45x = \$1,849/sh
- **P/E (FY27E):** EPS \$58.63 × 30x = \$1,759/sh; × 25x = \$1,466/sh

The cross-check brackets the DCF base case (~\$1,589) tightly: consensus FY27E EPS at 30x exit = \$1,759, within 7% of base DCF — independent triangulation confirms.

Entry-price range and position-sizing

ENTRY BAND	PRICE	VERSUS BASE \$1,589	IMPLIED 5Y IRR
Full position	≤ \$1,270	20% MoS	≥ +4.6% p.a. + asymmetry to bull
Starter position	\$1,270 – \$1,430	10–20% MoS	+2–5% p.a.
Watch / no-add	\$1,430 – \$1,650	0–10% MoS	<2% p.a. — current price sits here

ENTRY BAND

PRICE

VERSUS BASE \$1,589

IMPLIED 5Y IRR

Trim / exit

> \$2,180

37%+ premium

Negative IRR for forward 5y

POSITION-SIZING RECOMMENDATION

STARTER (small / 0.5–1.0% NAV). At \$1,648, current price sits just above the starter band and ~3% above the prob-weighted intrinsic. Quality-and-moat conviction is high (131 moat score); margin-of-safety wedge does not justify a full position. Recommended starter size at current price: ~\$5–10m per \$1bn of NAV, building only on retracement to $\leq$ \$1,430. Sell-side target average is \$2,230 [T2 — EODHD AnalystRatings] — Wall Street is paid to be at the upper end of our base–bull range.

Margin of safety — explicit calculation

- **Current price:** \$1,648.04 (mcap \$83.55bn $\div$ 50.697m shares)
- **Probability-weighted intrinsic:** \$1,632 ($0.50 \times \$1,589 + 0.25 \times \$2,726 + 0.25 \times \623)
- **Margin of safety:** $(1,632 - 1,648) \div 1,648 = -1.0\%$
- **Margin of safety to base case alone:** $(1,589 - 1,648) \div 1,648 = -3.6\%$
- **Required for full position (20% MoS):** $\leq \$1,270$

Valuation verdict: FAIR-leaning-EXPENSIVE. High-quality LatAm super-app trading at intrinsic value on consensus assumptions but with no margin of safety; the 49% Q1'26 re-acceleration is the call option, but it needs Q2 and Q3 2026 confirmation before sizing up from starter to full.

Indexes: Prob-weighted intrinsic \$1,632 vs spot \$1,648. MoS -1.0% . Bull-bear IRR spread 28.3pp. House WACC 13.5%.

§8 — INVESTMENT THESIS

Three structural anchors

ANCHOR 1 · THE INTEGRATED FIVE-ASSET STACK

MercadoLibre is not a single-vertical business. It is a marketplace + payments + credit + logistics + advertising platform serving the same customer base across 18 LatAm countries. Each asset reinforces the others — Mercado Pago settles marketplace transactions, Mercado Crédito extends working capital to marketplace sellers, Mercado Envíos delivers the goods that Mercado Pago settles, Mercado Ads monetises buyer attention, and the whole platform compounds on shared behavioural data. The combined moat is structural *[T1 — workstream_b §moat]*. Twenty-six years of compounded build-out is not replicable by a single competitor in a single quarter.

ANCHOR 2 · THE LATAM PENETRATION RUNWAY

E-commerce penetration in LatAm remains below 15% across the largest markets. The thesis is not that MercadoLibre captures a fixed slice of a fixed market — it is that the market itself doubles or triples while MercadoLibre holds or expands its share. The reference class for forward growth reads 38.4% (Sea Limited) at the comparable end and 49% (MELI Q1'26) at the recent end. The base rate supports the consensus 17–19% revenue CAGR; the call option is whether the Q1'26 re-acceleration is the new run rate or a one-quarter anomaly.

ANCHOR 3 · THE CASH-GENERATIVE MONETISATION ENGINE

TTM headline FCF of \$10.71bn is float-inflated; normalised operating FCF lands at \$3–4bn. Even on the lower number, MercadoLibre throws off cash at a quality and quantity that no LatAm e-commerce peer matches. The net cash position of \$12.1bn at Q1 2026 *[T1 — workstream_d §1]* is the structural balance-sheet defence against macro shocks in any single market.

Why now

The Q1 2026 +49% YoY revenue print is the call option

[T1 — workstream_d §2; 10-Q 2026-05-08]. It is the strongest YoY revenue print since Q2 2022. If

sustained through Q2 and Q3 2026, the bull-case scenario probability re-rates from 25% to ~40%, the probability-weighted intrinsic moves from \$1,632 to ~\$1,925, and the margin-of-safety wedge to the prevailing \$1,648 price opens up enough to justify a full position.

The CEO transition is structurally healthy [T2 — *workstream_c* §*founder-transition*]. Founder voluntarily handed off; internal CEO has a 9-year tenure including the Mercado Envíos build-out; bench is unusually long-tenured. The transition risk is real but compressed into the next four quarters of execution.

The price has corrected [T2 — *workstream_d* §1]. MELI is 38% below its 12-month high and 19% below its 200-DMA. The peak Mag-7-style multiple has compressed. We are buying into a corrected price, not a peak.

Why this position size

Starter, not full, because the AI disclosure does not yet justify the margin-expansion path embedded in the price. The DCF base case requires operating margin to expand from the FY25 11.08% to 18% over 10 years. The AI-claim driving the margin expansion is sourced from secondary aggregator synthesis — 87% AI customer-service resolution from klover.ai, 23.3% NIMAL from investing.com, 40% YoY fraud-cost decline from Gate 6 reference — not from primary management disclosure

[T1 — *cross_red_flags* §RF-DISC-006].

This is not a thesis-killer. It is a sizing constraint. Take a starter position today; demand AI-cohort primary anchoring as the gating condition for upsizing. If the Q2 2026 earnings call names a specific AI-driven cohort metric with attribution, upsize. If it does not, hold flat — and consider trimming if Q3 also defaults to scale-and-reinvestment.

The bull case — three concurrent operating conditions

The bull case is the call option, not the entry assumption. It pays off if three concurrent operating conditions hold through FY2026 and FY2027:

BULL CONDITION 1 · Q1 2026 +49% YOY IS NOT A ONE-QUARTER ANOMALY

If Q2 and Q3 2026 print at 35%+ YoY revenue growth on clean comps, the implied five-year revenue compound moves from the consensus 17–19% band toward the 25–30% band. The bull-case DCF scenario uses a 40/30/25/20/18% growth path with terminal operating margin reaching 22%, producing a **\$2,726/share** intrinsic *[T1 — [workstream_d §6\]](#)*. The +49% Q1 print materially raises probability mass on this path.

BULL CONDITION 2 · SZARFSZTEJN DELIVERS AI-DRIVEN MARGIN EXPANSION WITH ATTRIBUTION

The thesis-load-bearing AI claims — 23.3% Mercado Crédito NIMAL, 87% Mercado Pago AI customer-service resolution, 40% YoY fraud-cost decline — need primary management disclosure with cohort attribution. If Szarfsztein's Q2 2026 earnings call names a specific AI-driven cohort metric with attribution, the disclosure-quality flag retires, the AI-credit flywheel narrative anchors to T1 evidence, and the multiple defends.

BULL CONDITION 3 · SHOPEE RETREATS FROM LATAM

Sea Limited's board concludes that LatAm logistics-capex investment does not clear the SE-Asia opportunity cost. Shopee's LatAm subsidy programme tapers. MELI's take-rate normalises at the 20–22% band and Mercado Crédito grows uninterrupted.

If two of three bull conditions hold, the prob-weighted intrinsic moves from \$1,632 to approximately \$1,925. The starter position upsizes to a full position at the next entry-band window.

Sponsor framing — sell-side consensus and management's "deliberate reinvestment"

There is no sponsor in the traditional private-deal sense for this listed-equity screening paper. The closest analogues are (i) sell-side analyst consensus framing of the AI-

flywheel-driven margin-expansion thesis and (ii) management's own framing of "deliberate reinvestment" through FY2024 and FY2025.

On sell-side framing. Consensus models 17–19% revenue CAGR over five years with margin recovery to ~8.3% in ~3 years from ~6.9% trailing *[T2 — workstream_g §bull-framing]*. Our base case adopts a similar trajectory but with a slower margin recovery path (FY26 10.9% → FY30 16.0%) and a slightly heavier bear-case probability weighting. The sell-side bull framing is engaged, not silently inherited.

On management framing. Management invokes "deliberate reinvestment" to explain three consecutive years of operating-margin compression (FY24 12.66% → FY25 11.08% → Q1 2026 TTM 6.91%) *[T1 — workstream_d §2]*. The framing is plausible. But it has now been running for three full years; the market is 38% below the 12-month high precisely because investors are voting that the reinvestment is not buying the durability that management implies. We engage the framing explicitly: the partner discussion in §12 must ask how many more "deliberate reinvestment" years are tolerable before the multiple re-rates downward.

Falsifiability — five triggers that would break the thesis

1. Mercado Crédito 90+ day NPLs expand by more than 200 basis points in any single quarter without a quantitative explanation in the 10-Q segment data
[T1 — cross_red_flags §RF-ACCT-001].
2. Shopee LatAm executes a seller-subsidy programme that compresses MELI's marketplace take-rate from 20–22% to 17–18%
[T1 — workstream_b §war-game; workstream_d §6 bear].
3. BCB publishes a rule that caps platform-originated consumer credit at a sub-current-rate growth level *[T1 — workstream_d §9]*.
4. Szarfsztejn defaults to scale-and-reinvestment narrative through both Q2 and Q3 2026 calls without naming a specific AI cohort metric *[T1 — cross_red_flags §RF-DISC-006]*.
5. Operating margin compresses for a fourth consecutive year (FY26 < 6.91%)
[T1 — workstream_d §2; cross_red_flags §RF-ACCT-006].

Indexes: Thesis conviction HIGH. Sponsor framing engaged on both sell-side and management sides. Five binary falsifiability triggers.

BEAR CASE — THE OPERATIVE RISK THE STARTER SIZING CONCEDES

If Shopee LatAm escalates seller subsidies funded by SE-Asia cashflows, take-rate compresses from 20–22% to 17–18%. Simultaneously, BCB tightens fintech credit-origination rules, capping Mercado Crédito growth precisely when the AI-credit flywheel scales. Szarfsztein — five months in seat, with no operating track record in the CEO chair — fails to demonstrate margin expansion. The base-case probability mass compresses toward the bear scenario at **\$623/share**. The implied 5-year IRR from \$1,648 turns sharply negative at **–17.7% p.a.** *[T1 — workstream_d §6]*

This is a structurally credible scenario, not a tail outlier. The market is voting it some probability already: MELI is 38% below its 12-month high, the stock is 19% below its 200-DMA, and operating margin has compressed for three consecutive years.

The named risks

RISK 1 · SHOPEE LATAM SUBSIDY ESCALATION

Sea Limited's SE-Asia cashflows can fund an extended LatAm seller-subsidy programme that compresses MELI's take-rate from 20–22% to 17–18%. Central bear-case mechanism. **Mitigant:** Mercado Envíos owns first-party logistics across 18 countries; switching costs across Pago + Envíos + Crédito are high. **Trigger:** 100+ bps take-rate compression in any single quarter with no stated business reason.

[T1 — workstream_d §9; workstream_b §war-game]

RISK 6 · AI-NARRATIVE DISCLOSURE QUALITY

Thesis-load-bearing AI numbers (87% customer-service resolution, 23.3% NIMAL, 90% credit-book growth, 40% YoY fraud-cost decline) are sourced from secondary aggregator synthesis, not primary management disclosure. This is the high-severity flag in the 9-flag sweep. **Mitigant:** operating outcomes are real — Q4 2025 unit-shipping cost decline is primary-sourced; gross margins stable; 4-year revenue compound exceptional. **Trigger:** see falsifiability point 4.

[T1 — cross_red_flags §RF-DISC-006]

RISK 2 · BCB FINTECH CREDIT TIGHTENING

RISK 3 · SZARFSZTEJN EXECUTION GAP

BCB has a 12–24-month half-life on fintech credit-origination rules. Mercado Crédito growth is the operative fintech-margin contributor. **Mitigant:** Mercado Crédito is licensed; the BCB trajectory is gradual and consultative. **Trigger:** a BCB rule that caps platform-originated consumer credit at a sub-current-growth level.

[T1 — *workstream_d §9; workstream_b §regulatory*]

Galperín transitioned to Executive Chairman 1-Jan-2026; new CEO has not demonstrated through a competitive shock. **Mitigant:** nine-year MELI tenure including Mercado Envíos build-out; founder retains Executive Chair, capital allocation, AI sponsorship. **Trigger:** any quarter in which Szarfsztejn announces a major strategy departure without board endorsement.

[T1 — *workstream_d §9; workstream_c*]

RISK 4 · ARGENTINA FX TRANSLATION DRAG

BRL/ARS/MXN volatility erodes USD-translated revenue. IAS 29 hyperinflationary accounting adds noise. **Mitigant:** USD-reporting structure already absorbs FX; constant-currency disclosures in MD&A are the better signal. **Trigger:** >40% peso devaluation in a single quarter, or BCRA tightening capital repatriation.

[T1 — *workstream_d §9; workstream_e*]

RISK 7 · MARGIN COMPRESSION FOR A FOURTH YEAR

FY24 12.66% → FY25 11.08% → Q1 2026 TTM 6.91%. Three consecutive years of "deliberate reinvestment" framing without re-rate tracks the Schilit canonical pattern. **Mitigant:** management transparency; Mercado Crédito provisions and Envíos buildout identifiable. **Trigger:** FY26 op margin < 6.91% in clean (non-IAS-29) view.

[T1 — *workstream_d §2; cross_red_flags §RF-ACCT-006*]

RISK 5 · FLOAT-FCF DYNAMICS

TTM FCF \$10.71bn (33.7% margin) flatters; normalised operating FCF closer to \$3–4bn. **Mitigant:** deposit liability is visible on balance sheet; valuation is anchored on EV not raw FCF yield. **Trigger:** stripped operating FCF turns negative or contracts >25% YoY.

[T1 — *workstream_d §1*]

WHAT WE ARE NOT WORRIED ABOUT

The 86-flag sweep cleared 85 flags. No T1 finding in Regulatory or Legal. No Wirecard scope-limitation pair. No Tingo bank-confirmation pair. No going-concern modification, no late-filing pattern, no auditor change in last 3 years to smaller firm. No founder criminal record. No customer concentration. Mercado Envíos warehouses publicly verifiable across BR/MX/CL/CO. [T1 — *cross_red_flags §cleared*]

Indexes: 9 of 86 flags fired; 1 at high severity (RF-DISC-006). Bear scenario probability 25%. Bear IRR — 17.7% p.a.

Board composition

DIRECTOR	ROLE	INDEPENDENCE	BACKGROUND
Marcos Eduardo Galperín	Executive Chairman	No (Executive)	Founder; ex-CEO 1999–2025
Emiliano Calemzuk	Lead Independent Director	Yes	ex-Fox / News Corp
Susan Segal	Audit Committee	Yes	CEO, Americas Society / Council of the Americas
Richard Sanders	Independent	Yes	—
Alejandro Nicolás Aguzin	Independent	Yes	ex-JP Morgan, ex-HKEX CEO
Nicolás Roberto Galperín	Non-executive	No (related-party)	Family member (typically described as cousin)
Henrique Dubugras	Independent	Yes	Co-founder and co-CEO, Brex
Stelleo Tolda	Independent, Audit Committee	Yes (post-MELI)	ex-MELI EVP/COO and Commerce President; appointed Sept 2024
Martin Lawson	Independent	Yes	—

Independence ratio: 7 of 9 (~78%) *[T2 — workstream_c \$board]*. **Female representation:** 2 of 9. **Average board tenure:** ~9.3 years *[T3 — workstream_c]*.

Governance flags

- **Founder-Executive-Chair plus founder relative on board.** Related-party feature, not transaction. Disclosed. Doctrine canonical concern (family-controls-supplier pattern) does not apply. Low severity *[T1 — cross_red_flags \$RF-RP-001]*.

- **Lead Independent Director (Calemzuk)** partially mitigates the founder-Executive-Chair concentration [T1 — cross_red_flags §RF-GOV-002].
- **Long average board tenure (~9.3 years)**. Independence ratio is healthy; flag is on tenure depth. AC-chair-specific tenure not separately enumerated in current corpus — Gate 8 confirmation item [T1 — cross_red_flags §RF-GOV-001].
- **Zero board director with primary AI-research / AI-engineering identity**. Thesis leans on AI-flywheel; AI sits entirely inside management. Structural concern given entry justification [T2 — workstream_c §board].

Audit and disclosure

10-K filed 2026-02-25 for FY ending 2025-12-31 = 56 days; well within the 60-day large-accelerated-filer deadline [T1 — cross_red_flags §RF-GOV-012]. No going-concern language flagged. No auditor change in last 3 years identified. **Auditor identity is not in the current corpus** — presumed Big-4 (Deloitte historically) but must be confirmed before any move to upsizing [T1 — cross_red_flags §RF-AUD-001; substrate gap].

Compensation

CEO total comp approximately US\$4.3m with approximately 85% variable/stock-based and direct equity ownership approximately 0.0001% [T3 — workstream_c §compensation]. Founder Galperín's direct economic ownership and voting position is not crisply stated in the 2026 sources read in this pass [unverified — DEF 14A beneficial-ownership table not parsed]; based on prior public disclosures and Executive Chair posture, founder retains meaningful equity influence [inference]. Comp alignment index ~105 (around peer median). Listed-norm structure; no exotic features visible.

ESG posture

MELI operates as a licensed fintech in Brazil (BCB), Mexico (Ley Fintech), and Argentina (BCRA) — license stack is itself a regulatory-engagement signal [T1 — workstream_b §regulatory]. No publicly-named ESG dispute or activist-investor campaign identified in current corpus. Strategic posture toward LatAm consumer access is consistent with a constructive social ESG read [inference — workstream_a].

The structural ESG vulnerability is the cross-border LatAm regulatory perimeter: any single jurisdiction's politicisation of fintech credit is a material event. The Brazil concentration (~55% of revenue) is the named geography risk [T1 — cross_red_flags §RF-OPS-001].

Schilit-checklist screen

SCREEN	MELI READ	VERDICT
Dictatorial-CEO indicator	Galperín's voluntary step-down argues against; founder-Exec-Chair + family director is a watch overlay	Watch
Compensation skew	Heavily stock-linked, listed-norm; CEO total comp ~\$4.3m, 85% variable; ownership 0.0001%	Pass
Related-party transactions	Family director presence is disclosed feature, not flow	Pass (disclosed)
Auditor change signal	No auditor change in current filings; auditor identity itself not in corpus	Pass (substrate gap)
Kitchen-sinking write-off at transition	No write-off; 10-K filed 2026-02-25 reviewed clean	Pass

Integrity-gate verdict: clear-with-watch. Related-party board representation and long average tenure warrant partner attention but do not approach concern threshold.

Substrate gaps to close at Gate 8

- **Auditor identity** (presumed Big-4, likely Deloitte historically)

[T1 — cross_red_flags §RF-AUD-001]

- **Audit-fee trend** over three or more years
 - **PCAOB inspection findings** on the engagement
 - **AC-chair-specific tenure** within the long average board tenure of ~9.3 years
- [T1 — cross_red_flags §RF-GOV-001]
- **Galperín beneficial-ownership table** from the DEF 14A — voting position and share-pledge status [T1 — cross_red_flags §RF-CAP-004]

No going-concern language flagged in 10-K. No 8-K restatement filings identified in trailing 18-month primary_filings inventory.

Cleared governance flags

The cross-cutting red-flag sweep specifically cleared:

- No customer concentration (two-sided marketplace; millions of buyers)
- No founder criminal record for either Galperín or Szarfsztejn
- No fabricated credentials — Stanford MBA for both, UBA economics for Szarfsztejn (T2-corroborated)
- No pattern of attacking critics; no multiple sudden independent-director resignations
- No physical-asset verification gap — Mercado Envíos warehouses publicly verifiable from satellite imagery across BR/MX/CL/CO

Indexes: Independence ratio 78%. Average board tenure ~9.3y. Comp alignment ~105. Founder-courage marker ~110.

§11 — EXPERT REVIEW LOG

This pack was assembled from 13 healthy Workstream outputs across the listed-equity pipeline. Each workstream produced a v1 markdown one-pager plus structured JSON; each cleared verifier checks under the calibration-verification-pair discipline.

WORKSTREAM	OUTPUT	STATUS
A — Company History	19 events; 17 T1 / 2 T2	HEALTHY
A — AI Unit Economics	6 deployments; verdict HOLD; Tool User Gate FAIL	HEALTHY
B — Competitor Analysis	Helmer 7 Powers; moat index 125	HEALTHY
B — Voice of Customer	FLAG ONLY; floor not cleared	HEALTHY
C — People and Culture	CEO transition; bench depth	HEALTHY

WORKSTREAM	OUTPUT	STATUS
D — Valuation (bespoke)	DCF + reverse-DCF + scenarios; intrinsic \$1,632	HEALTHY
E — Primary Filings	Filings inventory; adversarial posture	HEALTHY
E — Earnings Transcripts	Q2-25 verbatim signal; bullish tone	HEALTHY
F — Reddit Intelligence	Community signal; 4 candidate themes flag-only	HEALTHY
F — CFPB Complaints	CFPB jurisdictional mismatch; reroute to LatAm	HEALTHY
G — Analyst Coverage	47 firms; Moderate Buy consensus	HEALTHY
H — Investor Sentiment	Mixed; no named short report	HEALTHY
Cross-cutting	9 fired / 7 watching / 27 substantively cleared / 43 categorical N/A	HEALTHY

Gate 7 health verdict: 13/13 HEALTHY

The verifier discipline produced four notable signal-shaping moves: (i) three thesis-load-bearing AI claims downgraded from primary to T4-secondary on re-read, firing the high-severity disclosure-quality red flag and shaping the conditional-trigger structure; (ii) margin trajectory fully T1-anchored to EODHD-parsed 10-K and 10-Q; (iii) CEO transition triangulated T1/T2/T3; (iv) Voice-of-Customer floor not cleared, so four candidate themes appear as colour only, not load-bearing.

§12 — IC DECISION RECORD

VERDICT

STARTER POSITION (CONDITIONAL)

Initial size 0.5–1.0% of NAV (~\$5–10m per \$1bn NAV) at spot (\$1,648). Scale to full position on retracement to the \$1,270–\$1,430 band. Hold at starter or trim if the conditional triggers fail to discharge by Q3 2026.

VIRTUAL IC DELIBERATION — 3 YES · 3 NO · 1 ABSTAIN

3

YES

Power-Law Hunter

Insurance Veteran

Operator

3

NO

Compounder

Macro Hawk

AI Skeptic

1

ABSTAIN

Activist — procedural

pending auditor identity +
DEF 14A beneficial ownership

Consensus verdict: STARTER (CONDITIONAL). Strategy and the Virtual IC consensus agree on STARTER. The dissent is internal to the Virtual IC panel.

Disagreement axis. The three NO votes (Compounder, Macro Hawk, AI Skeptic) argue the pack's own \$7 entry-price discipline places \$1,648 spot in the watch/no-add band — a full position belongs at \$1,270–\$1,430. The three YES votes (Power-Law Hunter, Insurance Veteran, Operator) read starter-at-spot as buying optionality on a binary AI-disclosure event at Q2/Q3 2026 earnings. Both readings are defensible from the same pack. The Activist abstain is procedural, pending auditor identity and Galperin DEF 14A beneficial-ownership data.

What the human IC should audit. Whether buying optionality on the AI-disclosure binary justifies a \$50k starter at –1.0% margin-of-safety, or whether Chronos price discipline argues for zero exposure until the pack's own \$1,270 trigger fires. The single highest-leverage question for management is whether Szarfsztein commits to publishing AI-driven cohort metrics with attribution at the Q2 2026 call — that disclosure resolves both the YES (validates moat) and NO (validates patience) sides on the same answer.

Minority view — serious dissent. Three of seven personas refuse to accept the –1.0% margin of safety as the price of optionality. The minority would re-vote YES if (i) price retraces into the \$1,270–\$1,430 starter band, or (ii) Q2 2026 anchors one AI cohort metric to primary disclosure with attribution. The human IC must address whether starter sizing materially changes Chronos's risk posture on MELI versus simply not yet being in the name.

Five questions the Human IC must ask management

1. **Will MercadoLibre publish one Mercado Crédito cohort default-rate metric — pre-AI baseline versus post-AI cohort at matched origination month — in the Q2 2026 earnings deck, or in the FY2026 10-K MD&A?** This is the single highest-leverage disclosure event in the next twelve months. The AI Skeptic's NO vote and the Power-Law Hunter's YES vote turn on the same answer.
2. **What is the current 90+ day NPL ratio on the Mercado Crédito book, disaggregated between seller credit and consumer credit, for the most recent quarter?** The walk-away trigger is set at +200 basis points in any single quarter without quantitative explanation. The IC cannot monitor a trigger it has never seen the baseline of. The Insurance Veteran's question, and it is binary.
3. **Who is the audit engagement partner, what is the three-year audit-fee trend, and what is Marcos Galperín's voting position and share-pledge status as of the most recent DEF 14A beneficial-ownership table?** Two substrate gaps in one question. The Activist abstains until these close. Both are routine disclosure items and the pack lists them as Gate 8 confirmation items that have not yet closed.
4. **Under what regulatory or macro condition would MercadoLibre commit to constant-currency primary reporting that is comparable across periods, and what is management's view on whether the IAS 29 hyperinflationary accounting treatment for Argentina is comparable across the FY24, FY25, and Q1'26 prints?** The Macro Hawk's vote turns on this. The pack itself names constant-currency reporting as a conditional trigger to upsize (\$0).
5. **Who is the named successor to Daniel Rabinovich as COO Product and Technology in the event of incapacity, and what is the timing of any public announcement of that succession?** The Operator's question. Twenty-six-year COO tenure across the entire product and technology stack is the single most load-bearing

operational concentration in the bench. The pack identifies it as medium key-person risk (\$5) but does not name a successor. A bench-depth question, not a hypothetical.

Conviction levels

- **Thesis conviction:** HIGH. Integrated five-asset stack is real, structurally moated, demonstrably compounding at the 4-year revenue CAGR of 42.2%. LatAm penetration runway is structural.
- **Price conviction:** FAIR-LEANING-EXPENSIVE. Probability-weighted intrinsic of \$1,632 against spot of \$1,648 implies –1.0% margin of safety. Current price embeds consensus 17–19% revenue CAGR at a developed-market discount rate.
- **Information sufficiency:** MEDIUM. AI-claim quality is the binding constraint. Three thesis-load-bearing AI metrics are not anchored to primary management disclosure. Auditor not named in current corpus. Galperín beneficial-ownership table not parsed.

Conditional triggers to upsize from starter to full

1. **AI cohort metrics anchored to primary disclosure.** Q2 or Q3 2026 earnings call (or FY2026 10-K MD&A) names at least one specific AI-driven cohort metric with attribution.
2. **Brazil regulatory clarity.** BCB publishes a rule drawing a stable perimeter around platform-originated consumer credit, or formally indicates moratorium for 12+ months.
3. **Argentina FX stability.** Sustained peso stabilisation through a full quarter, or MELI moves to constant-currency primary reporting comparable across periods.

Walk-away triggers (post-investment commitments)

The position is actively trimmed or exited if any of the following materialise post-entry:

- Mercado Crédito 90+ day NPLs expand by more than 200 basis points in any single quarter without quantitative management explanation.
- Marketplace take-rate compresses by more than 100 basis points in any single quarter.

- Operating margin compresses for a fourth consecutive year (FY26 < 6.91%).
- Szarfsztejn defaults to scale-and-reinvestment narrative through both Q2 and Q3 2026 calls without naming an AI cohort metric.
- Argentine peso devalues by more than 40% in any single quarter, or BCRA-led FX controls materially restrict capital repatriation.

Pre-mortem — three failure paths

FAILURE PATH 1 · SHOPEE LATAM SUBSIDY

Sea Limited executes a sustained LatAm seller-subsidy programme funded by SE-Asia cashflows. MELI marketplace take-rate compresses from 20–22% to 17–18%.

Kill metric: 100+ bps take-rate compression in any single reported quarter with no stated business reason.

FAILURE PATH 2 · BCB FINTECH CREDIT RULE

Brazil Central Bank publishes a rule capping platform-originated consumer credit growth at a sub-Mercado-Crédito-current-rate level. Mercado Crédito growth ceiling caps fintech-margin contribution.

Kill metric: published BCB rule explicit on platform-originated cap.

FAILURE PATH 3 · MARGIN COMPRESSION FOR A FOURTH YEAR

"Deliberate reinvestment" framing extends into FY26 without a re-rate. Operating margin TTM through FY26 < 6.91%. Market re-prices the multiple downward to mid-30s P/E.

Kill metric: FY26 operating margin reported below FY25 11.08% in a clean (non-IAS-29-distorted) view.

Dance-card next action

1. **Roger:** Sign off on starter position 0.5–1.0% NAV at spot, with explicit understanding that conditional triggers (1) and (2) must discharge before upsizing.

2. **Roger:** Commission a Gate 8.5 sub-workstream to anchor auditor identity, Galperín beneficial-ownership table from the DEF 14A, and Mercado Crédito 10-K segment data on NPLs and NIMAL. Substrate gap closure before Gate 9.
3. **Roger:** Diary Q2 2026 earnings call (expected mid-Q3 2026) as the AI-disclosure anchor event. Pre-build a one-page transcript-read template naming the specific phrases that would discharge the AI-cohort trigger.
4. **Roger:** Schedule a 30-minute partner review at the Q2 2026 earnings print to formally rerun the conditional-trigger decision and either upsize or hold.