

Strategy Verdict — MercadoLibre (MELI.US)

Spock v5 Strategy agent — eleven-framework synthesis with explicit pricing view. STARTER \$50,000 at -1.0% margin of safety. Generated 28 May 2026.

Strategy verdict — MELI.US

Headline

MercadoLibre is the integrated marketplace + payments + credit + logistics + advertising platform of Latin America, compounding revenue at 42.2% over the 2021–2025 window and re-accelerating to +49% YoY in Q1 2026 [claim_002, claim_003]. At spot \$1,648 the probability-weighted intrinsic is \$1,632, implying –1.0% margin of safety [claim_001]. The verdict is **STARTER** at ~\$50k, not full. The STARTER at a negative margin of safety is deliberate and defensible: bull-case upside to \$2,726/share is real and asymmetric if the Q1 2026 re-acceleration sustains, the integrated five-asset moat (Helmer 7 Powers weighted 115.7) is structural rather than AI-exceptionalist, and the Chronos LatAm-fintech edge underwrites the patience cost of a small real-money seat while the AI-disclosure trigger discharges [claim_004, claim_005]. Full position only at ≤\$1,270 (~20% MoS) [claim_017]. The thesis is high; the price is fair-leaning-expensive; the information sufficiency is medium. Take the seat, demand the disclosure, scale on triggers.

The eleven frameworks

Power law — Partial pass. Bull \$2,726 (+10.6% 5y IRR) vs bear \$623 (–17.7%); 28.3pp spread narrowly below the 30pp DCF Standard floor, which reflects MELI's \$31.8bn revenue base, not a thesis weakness [claim_015].

CLTV:CAC — Not applied; no cohort disclosure in primary filings. TTM ROE 31.26% and four-year revenue CAGR 42.2% argue aggregate returns are strong, but Michael's 9× bar cannot be cleanly tested. Same disclosure gap as RF-DISC-006 [claim_002].

Helmer 7 Powers — Wide-moat candidate. Scale economies 135, network economies 135, switching costs 130 all load on primary 10-K. Cornered resource 95: 26 years of LatAm behavioural data is replicable with time [claim_004, claim_005].

AI Maturity — Level 2 leaning 3. Zero of six deployments grade B-or-better; Tool User Gate fails fail-closed; three bullshit flags fire. Strongest deployment (Mercado Ads) at B-. Structural concern vs Nubank's Level 3+ AI-native credit underwriting [claim_006, claim_007].

Sizing — Starter \$50,000. Punt-tier dollar size, below normal \$250k pending trigger discharge. Consistent with CHYM precedent set 27 May 2026 IC. Upsize at next entry-band window if Q2 2026 Szarfsztein call names an AI-cohort metric [claim_022].

Scenarios — Base 50% / Bull 25% / Bear 25%. Probability-weighted intrinsic \$1,632. Bear is load-bearing and structurally credible; the starter sizing surfaces it explicitly rather than glossing it [claim_015].

Listed-v-private divergence — Pass. Integrated five-asset LatAm platform across 18 countries does not duplicate the Chronos private book; the geography is a diversified overlay [claim_013].

Geographic bias — Explicitly defended. International over SA criterion satisfied. Brazil >50% fires RF-OPS-001 medium; defended as regional-champion exposure with US-listed USD-reporting absorbing FX [claim_013].

Trim discipline — No trim contemplated. Position not yet entered; Willem's rule does not engage. Forward trim zone above \$2,180 documented as partner consideration only [claim_017].

Disagreement with Virtual IC — Pending. Virtual IC running in parallel; to be re-evaluated post-merge.

Pricing — Fair-leaning-expensive. Five-method triangulation brackets the base case tightly: P/E forward 31.95x fair vs Amazon 34.9x; EV/EBITDA 23.84x stretched vs Sea 20.2x; EV/Revenue 2.84x fair; DCF base \$1,589; FY27E EPS \$58.63 $\times$ 30x = \$1,759 (within 7% of DCF base) [claim_014]. Reverse-DCF embeds 18% revenue CAGR at 13.5% WACC [claim_016]. Margin of safety –1.0%.

Bear case

Structurally credible at \$623/share (–17.7% 5y IRR from spot); the market is already pricing some of it (stock 38% below 12-month high, 19% below 200-DMA) [claim_024]. Three concurrent mechanisms: (1) Shopee LatAm seller-subsidy escalation funded by Sea’s SE-Asia cashflows compresses take-rate from 20–22% to 17–18% [claim_011]; (2) BCB tightens fintech credit-origination rules with 12–24 month half-life, capping Mercado Crédito growth precisely when the AI-credit flywheel scales [claim_012]; (3) Szarfsztejn — five months in seat, first non-founder CEO — fails to demonstrate margin expansion through both Q2 and Q3 2026, defaulting to a fourth year of “deliberate reinvestment” framing [claim_010]. Operating margin has already compressed FY24 12.66% → FY25 11.08% → Q1 2026 TTM 6.91%. The bear is the active risk the starter sizing concedes.

Conditional triggers to upsize

1. **AI cohort metrics anchored to primary disclosure.** Szarfsztejn Q2 or Q3 2026 earnings call names a specific AI-driven cohort metric with attribution. Discharges RF-DISC-006 and re-rates bull probability 25% → 40%, moving probability-weighted intrinsic to ~\$1,925 [claim_006].
2. **Brazil regulatory clarity.** BCB publishes a rule drawing a stable perimeter around platform-originated consumer credit, or formally indicates a 12+ month moratorium on further tightening [claim_012].
3. **Q1 2026 +49% YoY confirmed as new run-rate.** Q2 and Q3 2026 print 35%+ YoY on clean comps [claim_003].
4. **Argentina FX stability or constant-currency reporting.** Removes IAS 29 hyperinflationary noise from headline GAAP.

Disagreement with Virtual IC

Placeholder — running in parallel. Strategy verdict is STARTER. Post-merge, the human IC should audit which axis dominates if disagreement surfaces: information-sufficiency objection (sizing-only disagreement) or bear-case probability weighting (thesis-level disagreement). Either is non-trivial and should be surfaced explicitly to partners, not smoothed.

Spock v4 listed-eq pipeline · MELI.US · generated 28 May 2026.