

PREVIEW DISCLOSURE

This is a Gate 6.5 high-conviction Strategy preview, not a final IC paper. Workstreams B (voice-of-customer / consumer signal), C (people & culture), D (full scenario derivation tables) and the Virtual IC have not been run. Public-source baseline + Strategy-agent discipline only. Every load-bearing claim cites a public source or is tagged [unverified]. Indexing convention: 100 = median, range ~50–150.

SPOCK V5 — STRATEGY FRAMEWORK PREVIEW

Verisk Analytics: the quietest moat in US insurance

A 60-year ISO franchise, a near-monopoly on US property & casualty loss data, and an AI surface area that is real but undisclosed. The question is not whether it is a good business — it is whether the price still allows for a Chronos return.

VRSK.US · INFORMATION SERVICES · PROPRIETARY DATA HOLDER · 28 MAY 2026

STRATEGY VERDICT

Build a starter position — accumulate on weakness, do not chase

Verisk is a textbook Helmer "cornered resource" with insurance-carrier switching costs and a regulated-data network effect that compounds quietly. The bull case requires AI-product monetisation on top of the legacy ISO franchise; the bear case is a low-growth utility re-rate. At consensus forward multiples the equity offers a steady compounder return, not a Chronos asymmetry. Size as starter; let the price come to you.

RECOMMENDATION

Starter (build on weakness)

INITIAL SIZE

USD 50,000

SIZING BUCKET

Punt → starter ladder

VALUATION READ

Fair-to-rich; ~10–15% MoS req'd

Verisk Analytics owns the data layer underneath the US property & casualty (P&C) insurance industry. Its ISO statistical-agent franchise, codified in state insurance regulation since the 1970s, is a privileged information utility that competitors cannot rebuild because they cannot collect the underlying carrier-pooled loss data [verisk.com/about/iso]. The post-2023 simplification — divesting Wood Mackenzie (energy) and the financial-services unit — has left a pure-play insurance-data company with high-teens EBITDA margins and consistent net revenue retention [sec.gov 10-K filings]. The strategic question for Chronos is whether the franchise plus a credible (but undisclosed in scale) AI-product layer justifies paying mid-thirties P/E in a market where the equivalent insurance-domain exposure is already partially owned through Hollard, OUTsurance and the private InsurTech overlay.

The eleven Strategy frameworks

1. Power law

NOT ASYMMETRIC

The Spock Power-law screen asks: at a five-year horizon, can the equity plausibly return 3x or more on the base case, with a long-tail bull beyond that? Verisk fails this on the upside. The 2019–2024 revenue CAGR on the simplified insurance-only perimeter is roughly high-single-digit organic [investor.verisk.com — [quarterly results](#)], with margin expansion adding another 100–200 bps per year of EPS growth and buybacks contributing the residual. That is a 12–15% total-shareholder-return business on the base, not a 25%+ compounder. The upside tail is capped because Verisk does not have a credible adjacent-market expansion story since the Wood Mackenzie divestiture [[verisk press, 2023](#)]. The downside is shallow — the franchise is hard to break — but the equity is engineered for steady compounding, not asymmetric payoff.

Power-law verdict: **fail**. Treat as a quality-compounder ballast position, not a thesis bet.

2. CLTV : CAC and unit economics **BEST-IN-CLASS**

Verisk's customer base is the US P&C insurance industry — a finite, slow-moving universe of carriers and reinsurers that embed Verisk products into underwriting and claims workflows. Subscription revenue is approximately 80%+ of total revenue per disclosure across recent 10-Q filings [[sec.gov 10-Q](#)], with multi-year contracts and high renewal rates. Management has historically guided to organic constant-currency revenue growth in the 6–9% range; Verisk has not, to my knowledge, disclosed a formal Net Revenue Retention number in the SaaS sense [unverified — NRR not publicly disclosed in standard form], but renewal rates in the insurance-services segment have been described as "in the high 90s" in past investor communications [unverified — investor day characterisation]. CAC is structurally low because the customer set is enumerable and the sales motion is enterprise-renewal, not acquisition. CLTV:CAC verdict: **excellent**, but reflected in the multiple. Unit economics are the strongest single argument for the name.

3. Helmer 7 Powers **3-OF-7 STRONG**

Three powers screen as clearly present. **Cornered Resource:** the ISO statistical-agent designation and the historical pooled P&C loss database — built since the 1970s under state insurance department arrangements — is genuinely irreplacable; a new entrant cannot collect the underlying data because carriers will not share with a non-incumbent agent [[verisk.com/about/iso](#)]. **Switching Costs:** Verisk products (ISO ERC, ClaimSearch, Xactware, AIR/Extreme Event Solutions) are embedded into carrier underwriting and claims workflows, regulatory filings, and rate manuals; rip-and-replace cycles are multi-year and politically costly inside carriers [[verisk.com/insurance](#)]. **Scale Economies:** the marginal cost of serving an additional carrier on a built data asset is near-zero; gross margins reflect this. The four other powers are weaker but not absent. **Network economies:** partial — Verisk's contributory databases (ClaimSearch in particular) get stronger as more carriers contribute, but the network is closed-loop within the existing customer set, not a platform-economy flywheel that attracts new entrants. **Branding:** present but secondary

— "Verisk" and "ISO" carry weight inside the insurance industry but the brand is not a consumer-facing moat. **Counter-positioning:** weak — incumbents like LexisNexis Risk Solutions and TransUnion compete in adjacent risk-data verticals and would respond symmetrically to a price-disrupting move. **Process power:** weak — the moat is the data and the regulatory designation, not a hard-to-copy operating cadence. Three structural powers stacked on one customer set is unusual and is the core reason the franchise prints high-30s EBITDA margins. The risk inside Helmer is that the three strong powers are static — they protect what is, not what could be. The AI-product layer is where Verisk would need to add a fourth power (network economies in the platform sense around AI training data), and that has not yet shown up.

4. AI Maturity Level 1–5 LEVEL 2 – NARRATIVE_ONLY

Apply Spock doctrine v1.1: Verisk is a **Proprietary Data Holder** archetype. The doctrinal expectation is high — a company sitting on the most valuable proprietary insurance-loss dataset in the United States should be at AI Maturity 3 or 4 (production AI products with disclosed customer adoption metrics). What Verisk actually discloses falls into the **narrative_only** bucket on the disclosure_quality_axis: press releases describe AI-enabled fraud detection, generative-AI underwriting assistants, and predictive analytics overlays on ClaimSearch and Xactware [[verisk.com](https://www.verisk.com) — product pages], but neither the 10-K nor recent earnings transcripts contain audited AI revenue, AI ARR, or AI customer-count metrics [unverified — searched 2024 and 2025 10-K/Q disclosures; nothing audited_quantitative found]. This is the doctrinal red flag for Proprietary Data Holders: if the AI story is real, it should show up in the numbers. Verdict: **Level 2 of 5**, with upside optionality if management starts disclosing a quantitative AI line in 2026. The doctrine v1.1 critique sharpens further: for a Proprietary Data Holder, "AI capability" without disclosed customer attach and revenue is indistinguishable from marketing — the data asset is necessary but not sufficient. Verisk's peers in the disclosure-quality comparison are uneven: S&P Global discloses Kensho-related AI capabilities with limited quantitative anchoring; MSCI has begun disclosing AI-augmented index product detail; Moody's has been more explicit about generative-AI product attach in its 2024–25 communications [unverified — peer disclosure characterisations to be re-anchored at IC stage]. Against that benchmark Verisk is behind on disclosure cadence, which matters

less for the franchise itself and more for the multiple — equity owners are being asked to take AI optionality on trust. The Strategy view: until Verisk moves to at least management_disclosed_quantitative (a CEO or CFO stating an AI revenue figure on an earnings call with sourcing), assume Level 2 and price the optionality at zero.

5. Sizing bucket PUNT → STARTER

The Spock sizing ladder is high-conviction \$1M / normal \$250k / punt \$50k / starter. Verisk does not earn high-conviction on this preview because Power-law fails and AI disclosure is narrative-only. Normal \$250k is the natural slot for a quality compounder with a defensible moat, but the entry multiple is the binding constraint — paying 32–35x forward earnings [unverified — exact consensus FY26 P/E to be re-pulled at IC stage] for a 12–15% TSR business eats most of the margin of safety. Recommendation: **\$50,000 starter** with a written escalation rule — add 50% on a 15% drawdown without thesis break, add another 100% (to ~\$200k total) on a 25% drawdown or on the first audited_quantitative AI revenue disclosure. Sizing bucket verdict: **punt → starter**, not a normal slot.

6. Bull / Base / Bear scenarios SKEWED FLAT

Probabilities (Strategy preview, not Workstream D derivation):

SCENARIO	PROBABILITY	5YR CAGR (PRICE)	DRIVER
Bull	25%	~18–22%	AI products land at scale; ISO-style monetisation on AI-augmented claims and underwriting; carrier-side AI budget flows to incumbent data vendor; modest re-rating to 38–40x forward.
Base	55%	~9–12%	Steady 7–8% organic revenue, ~100bps margin expansion, buybacks; multiple flat at ~32x; total return ≈ FCF yield + growth.

SCENARIO	PROBABILITY	5YR CAGR (PRICE)	DRIVER
Bear	20%	~ -2 to +3%	AI commoditises basic underwriting analytics; insurance customer consolidation reduces seats; multiple compresses to 22–25x; flat-to-down absolute.

Probability-weighted ~10–12% 5yr CAGR. [unverified — probabilities are Strategy-preview judgement, not Workstream D derived] . This is the math behind the Power-law fail: even the bull case does not produce a 3x.

7. Listed-vs-private divergence (Chronos cross-read)

GENUINE ADDITIVE EXPOSURE

Chronos has insurance-domain exposure already through Hollard (Nic Kohler), OUTsurance (Willem Roos), Discovery, and the regulated TIC + InsurTech private overlay. Those exposures are **carrier-side, primarily South African and Australian**. Verisk is **data-infrastructure-side, US-centric**. The cross-read: Verisk adds an exposure the private book structurally cannot replicate — you cannot buy a US P&C statistical-agent franchise; it is regulated and incumbent [verisk.com/about/iso]. The divergence is genuine and the partners' domain expertise is a real edge in evaluating Verisk's customer-side risk (carrier consolidation, MGA disintermediation, AI underwriting in-house insourcing). Verdict: **additive**. The Chronos-specific argument for owning Verisk is stronger than the generic information-services argument; we can underwrite the customer risk better than a generalist.

8. Geographic bias

US-CONCENTRATED — DEFENDED

Approximately 80%+ of revenue is United States [unverified — exact geographic mix to be re-pulled from latest 10-K segment disclosure] , with the balance UK, Europe, and modest Asia-Pacific. The Chronos book is structurally under-weight US insurance-data infrastructure; adding US exposure here is a portfolio diversifier, not a concentration. The geographic-bias concern at portfolio level (over-USD, over-US listed equity in the Chronos book) is real but Verisk's USD revenue base is

naturally hedged to USD liabilities and the franchise's regulatory protection is specifically a US-state-regulator phenomenon — moving offshore would dilute the moat. Defend, do not downgrade. Verdict: **acceptable US concentration**, with a watch on whether the broader Chronos listed book has crossed an internal USD concentration threshold [unverified — internal portfolio limit not checked in this preview] .

9. Trim discipline **NO POSITION YET — N/A**

Chronos does not currently hold Verisk to my knowledge [unverified — portfolio system not queried in this preview] . Trim discipline therefore not engaged. Forward rule to be written into the entry note: **trim 25% if forward P/E exceeds 40x without an accompanying audited AI revenue disclosure; trim 50% if forward P/E exceeds 45x in any scenario; review position at every 15% rally without earnings revision**. Encoding the trim rule at entry is the discipline; the framework's purpose here is to prevent the "quality forever" failure mode where a fair-to-rich entry compounds into an expensive holding nobody re-underwrites.

10. Disagreement with Virtual IC **NOT RUN — PLACEHOLDER**

Virtual IC has not been run for this preview (Gate 9 not executed). Placeholder. The dimensions on which the Strategy view would expect challenge: (a) Power-law fail — a competing view would argue Verisk's optionality on AI monetisation justifies a higher base-case probability than 25%; (b) starter sizing — a competing view would argue the franchise quality merits a \$250k normal slot even at the multiple; (c) bear-case weight — a competing view might argue 20% understates AI commoditisation risk to legacy analytics. The Strategy agent's position to be defended in Gate 9: **quality is real, but quality is also priced; do not let franchise-love override the math**.

11. Pricing / valuation view **FAIR-TO-RICH**

Multi-method read, peer set: S&P Global (SPGI), Moody's (MCO), MSCI (MSCI), FactSet (FDS), Morningstar (MORN). All disclose forward P/E and

EV/EBITDA via standard data providers [unverified — current consensus to be pulled at IC

stage; figures below are best-public-baseline as of recent quarters] .

METHOD	VRSK	PEER MEDIAN	READ
Forward P/E (FY26E)	~32–35x [unv]	~28–32x [unv]	10–15% premium to peers
EV/EBITDA (FY26E)	~22–24x [unv]	~19–22x [unv]	Modest premium
EV / organic revenue growth	~3.5–4.0x growth- adj [unv]	~3.0–3.5x [unv]	Slight premium
FCF yield (FY26E)	~3.0–3.5% [unv]	~3.5–4.0% [unv]	Below-peer FCF yield
Implied IRR (Base scenario, 5yr)	~9–12%	—	Acceptable, not Chronos- target
Margin of safety to fair value	~ -10 to -15% [unv]	—	Currently negative; wait for drawdown

Verdict: **fair-to-rich**. Verisk trades at a modest premium to its information-services peer set, which is defensible on franchise quality but does not leave a Chronos margin of safety. Required entry: 10–15% drawdown from current levels, or first audited_quantitative AI revenue disclosure, whichever comes first. [unverified — all multiples carry "unv" pending IC-stage consensus pull from a primary data provider] .

Headline thesis

Verisk is the quietest moat in US insurance — a 60-year ISO statistical-agent franchise that competitors structurally cannot rebuild, embedded into carrier workflows under state insurance regulation, and now simplified post-Wood Mackenzie into a pure-play insurance-data infrastructure name. The unit economics are best-in-class: ~80% subscription, high-90s renewal characterisation, near-zero marginal cost on a built data asset. Three of

Helmer's seven powers stack on a single customer set, which is rare. The Chronos-specific argument is unusually strong: our carrier-side domain expertise (Hollard, OUTsurance, Discovery, InsurTech overlay) lets us underwrite Verisk's customer-side risks (carrier consolidation, AI insourcing, MGA disintermediation) better than a generalist owner. The break in the thesis is price and AI disclosure. Verisk trades at a modest premium to its information-services peer set, which leaves no Chronos-grade margin of safety on entry; and the AI story — which is doctrinally critical for a Proprietary Data Holder — is narrative_only, not audited_quantitative. Build a starter; let the price come; escalate on the first audited AI revenue disclosure or on a 15% drawdown without thesis break.

BEAR CASE — READ FIRST

Two compounding bear drivers, each individually live, jointly devastating. **AI commoditisation of legacy analytics:** if foundation-model-class AI lets a Tier-1 carrier or a well-capitalised insurtech replicate ~70% of Verisk's predictive-analytics value internally using public + carrier-internal data, the franchise narrows to the irreplaceable ISO loss-pool layer — still defensible, but a thinner, slower-growth product. **US insurance carrier consolidation:** the US P&C industry is mid-consolidation; fewer, larger carriers means fewer seats, harder pricing conversations, and a higher proportion of revenue concentrated in customers with internal data-science teams. Combine the two and the multiple compresses from 32x to the low-20s; combine that with high-single-digit revenue growth and the equity returns flat-to-down for three years. The bear case is not catastrophe — it is utility-isation. The thesis-killer is the first audited_quantitative AI revenue line that fails to materialise alongside continued carrier consolidation. [unverified probabilities

— 20% on the explicit bear, but partner challenge expected on this weight] .

Pricing summary — multi-method

The pricing read across five methods is internally consistent: Verisk is fair-to-rich, with no current Chronos-grade margin of safety. The cleanest single read is the FCF yield: at $\sim 3.0\text{--}3.5\%$ the equity prices the franchise as a high-quality bond-substitute, which is defensible at current US 10yr levels but vulnerable to either rate normalisation or earnings disappointment. The information-services peer set (SPGI, MCO, MSCI, FDS, MORN) is the right comparator because all five share the proprietary-data-on-subscription business model; cross-sector comparators (e.g. industrial software) are misleading because the regulatory-utility character of Verisk's ISO franchise has no equivalent in horizontal software. The pricing watch: any move below $\sim 28\times$ forward P/E without an earnings cut is a Chronos buy signal; any move above $\sim 38\times$ forward P/E without an AI revenue disclosure is a Chronos trim signal. [unverified — exact multiple

bands to be re-anchored at IC stage from a primary data provider] .

Two cross-checks reinforce the read. The first is the **EV / organic-growth** walk: Verisk's $\sim 7\text{--}8\%$ organic revenue growth at $\sim 22\text{--}24\times$ EV/EBITDA implies the market is paying for either margin expansion (plausible, but bounded — high-30s EBITDA margin leaves limited runway) or for AI-driven re-acceleration (not yet evidenced). The second is the **implied IRR** walk: at the current FCF yield plus consensus growth, the equity offers $\sim 9\text{--}12\%$ on the Base scenario, which is below the Chronos $15\%+$ target for normal-sized positions but acceptable for a starter ballast slot. The math is internally consistent across methods — the conclusion is not that any single multiple is mispriced but that the entire valuation stack assumes the AI optionality will be delivered, which is the precise gap the AI Maturity framework flagged.

Insurance-domain cross-read — Chronos-specific

This is the section that distinguishes a Chronos read of Verisk from a generalist information-services read. Chronos partners carry direct, current insurance-

carrier operating experience: Nic Kohler (Hollard), Willem Roos (OUTsurance), Discovery, and the InsurTech overlay thesis through regulated TIC + adjacent private holdings. That domain expertise has three load-bearing implications for the Verisk thesis.

First, customer-side risk reads better. A generalist owner cannot easily judge whether US P&C carriers will insource Verisk's analytics into their own AI/ML stacks. Chronos can — because the partners have been on the carrier side of the same conversation in different jurisdictions. Initial read: large carriers will partially insource the AI-overlay layer but will not insource the ISO loss-pool layer because they cannot recreate the data [verisk.com/about/iso]. That is the same intuition any senior insurance operator would have, but priced into the Chronos thesis it tightens the bear case.

Second, the listed-vs-private divergence is genuine. The Chronos private book has no US P&C data-infrastructure exposure and structurally cannot acquire it — the ISO franchise is regulated and incumbent. Verisk is therefore not redundant against existing holdings; it is additive in a way few US-listed insurance-domain names would be.

Third, partner-channel diligence is available. Both Nic Kohler and Willem Roos are positioned to provide a short, structured customer-side read on Verisk's product portfolio — ClaimSearch, Xactware, AIR/Extreme Event Solutions, ISO ERC. That is a Gate 7 Workstream-B substitute that a generalist screener cannot run. Recommended pre-IC: a 30-minute structured call with each of them, with a fixed five-question template covering switching costs, AI-insourcing risk, MGA disintermediation, pricing-power durability, and the carrier view of Verisk pricing increases over the last three renewal cycles.

Three escalation triggers

1. **First audited_quantitative AI revenue disclosure.** If Verisk discloses a dedicated AI revenue line (ARR, % of revenue, customer count, or attach rate) in a 10-K, 10-Q, or investor day with audited figures — escalate from starter to normal sizing (\$50k → \$200k), conditional on the disclosed AI

line being > 5% of revenue and growing faster than the base business. This is the single highest-signal event for moving Verisk out of "narrative_only" maturity.

2. **15% drawdown without thesis break.** If the equity drops 15% from current levels while organic revenue growth, renewal rates, and EBITDA margin remain within guided bands — add 50% to position (build from \$50k to ~\$75k). This is the price-comes-to-us escalation. If the drawdown is accompanied by a thesis break (customer loss, AI insourcing announcement from a top-5 carrier, multiple-quarter organic deceleration) — do not add; reassess.
3. **Carrier consolidation event.** If a top-10 US P&C carrier announces a strategic decision to insource Verisk-equivalent analytics, or if a top-5 carrier merges with another top-10 — trigger a pre-IC review within 30 days. This is the bear-case canary. Outcome: either the bear weight increases to 30%+ and we trim/exit, or the carrier-side read confirms the ISO franchise is unaffected and we hold.